



SuperDataScience

COMPUTER VISION




<div>Part 1</div> <div>Face detection</div>	<div>Part 2</div> <div>Face recognition</div>	<div>Part 3</div> <div>Object tracking</div>	<div>Part 4</div> <div>Neural networks for image classification</div>	<div>Part 5</div> <div>Convolutional neural networks for image classification</div>
<div>Part 6</div> <div>Transfer learning and fine tuning</div>	<div>Part 7</div> <div>Neural networks for classification of emotions</div>	<div>Part 8</div> <div>Autoencoders</div>	<div>Part 9</div> <div>Object detection with YOLO</div>	<div>Part 10</div> <div>Recognition of gestures and actions</div>
<div>Part 11</div> <div>Deep dream</div>	<div>Part 12</div> <div>Style transfer</div>	<div>Part 13</div> <div>GANs (Generative Adversarial Networks)</div>	<div>Part 14</div> <div>Image segmentation</div>	

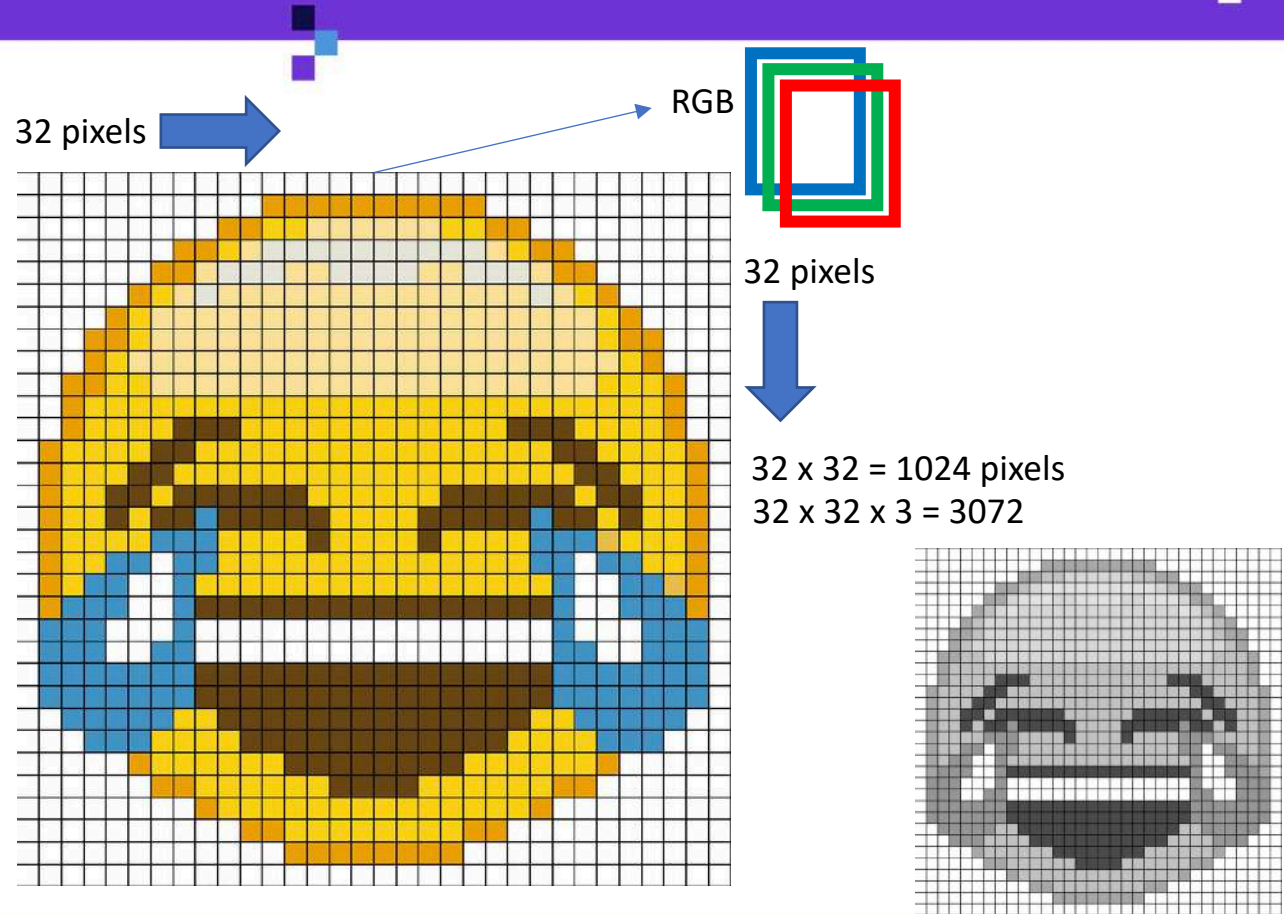
- Basic python programming
- Only the basic intuition (except for neural networks)
- Libraries
 - OpenCV
 - Dlib
 - TensorFlow
 - Darknet
 - Caffe framework
- Homeworks

PLAN OF ATTACK – FACE DETECTION

1. Face detection with Haarcascade and OpenCV
2. Face detection with HOG and Dlib
3. Face detection with CNN and Dlib
4. Face detection using webcam



PIXELS



CASCADE CLASSIFIER




Faces



Not faces

AdaBoost
Training

Feature
selection




Apply to each "sub-window"

CASCADE CLASSIFIER




Sum of white pixels – sum
of black pixels




2 3 5 6
8 9 2 1


CASCADE CLASSIFIER







CASCADE CLASSIFIER



HOG – HISTOGRAMS OF ORIENTED GRADIENTS



Original Image : 720 x 475



2	3	4	4	3	4	2	2
5	11	17	13	7	9	3	4
11	21	23	27	22	17	4	6
23	99	165	135	85	32	26	2
91	155	133	136	144	152	57	28
96	196	76	38	26	60	170	51
165	60	60	27	77	85	43	136
71	13	34	23	108	27	48	110

Gradient Magnitude

80	36	5	10	0	64	90	73
37	9	9	179	78	27	169	166
87	136	173	39	102	163	152	176
76	13	1	168	159	22	125	143
120	70	14	150	145	144	145	143
58	86	119	98	100	101	133	113
30	65	157	75	78	165	145	124
11	170	91	4	110	17	133	110



Gradient Direction

Derivative allows to measure the rate of change
(zero derivative, small derivative and high derivative)


Gradient vector (direction that the values increase)

Source of images: <https://www.learnopencv.com/histogram-of-oriented-gradients/>

HOG – HISTOGRAMS OF ORIENTED GRADIENTS




HOG – HISTOGRAMS OF ORIENTED GRADIENTS



PLAN OF ATTACK – FACE RECOGNITION

1. Face recognition with LBPH and OpenCV
2. Face recognition with Dlib, CNN and distance calculation
3. Face recognition using the webcam




Pred: 10
Exp: 10



LBPH (LOCAL BINARY PATTERNS HISTOGRAMS)

12	15	18
5	8	3
8	1	2

If ≥ 8 : 1
If < 8 : 0



1	1	1
0	8	0
1	0	0

Binary = 11100010

Decimal = 226





42	55	48
35	38	33
38	30	32

Binary = 11100010


LBPH (LOCAL BINARY PATTERNS HISTOGRAMS)

Edge

Corner





Source: https://bytefish.de/blog/local_binary_patterns/



LBPH (LOCAL BINARY PATTERNS HISTOGRAMS)

1. Radius
2. Neighbors
3. grid_x and grid_y
4. Threshold



12	15	18
5	8	3
8	1	2



Source: https://en.wikipedia.org/wiki/Local_binary_patterns

PLAN OF ATTACK – OBJECT TRACKING

1. Object tracking vs. Object detection
2. KCF (Kernel Correlation Filters) algorithm
3. CSRT (Discriminative Correlation Filter with Channel and Spatial Reliability) algorithm
4. KCF and CSRT implementation

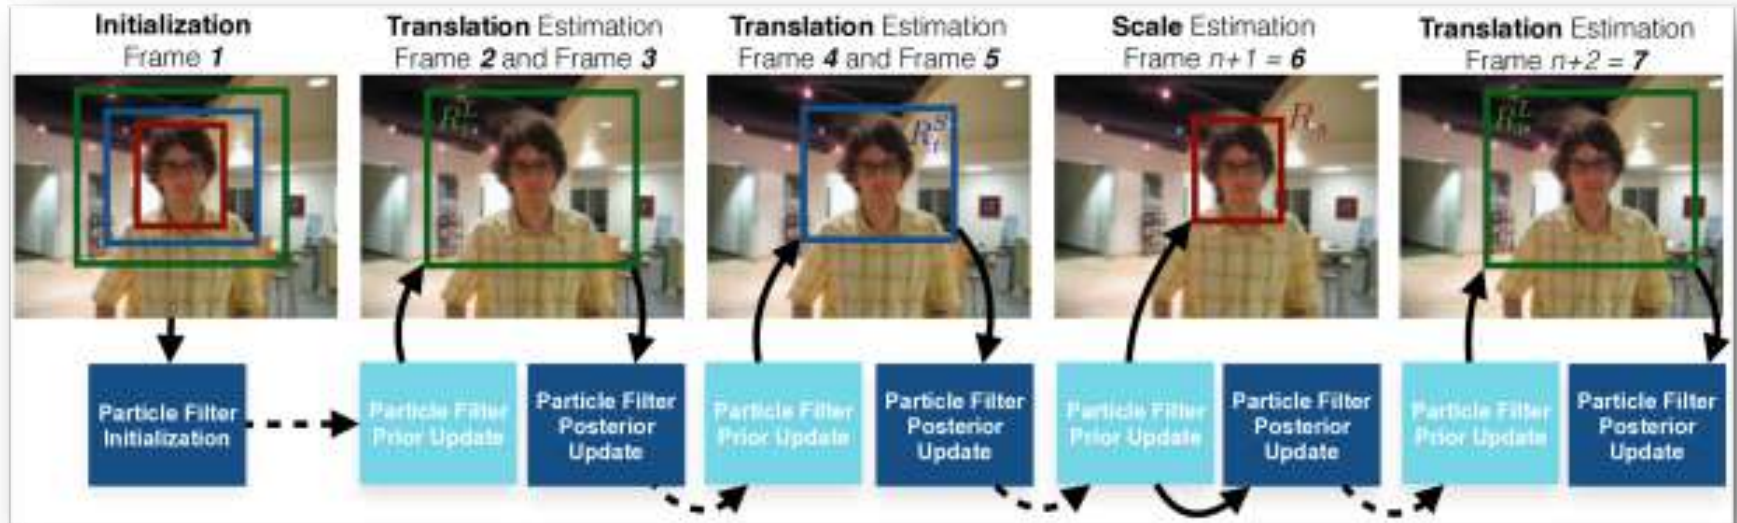


OBJECT TRACKING vs. OBJECT DETECTION

- Object tracking vs. Object detection




KCF (KERNAL CORRELATION FILTERS)



Source: <https://www.groundai.com/project/enkcf-ensemble-of-kernelized-correlation-filters-for-high-speed-object-tracking/1>

CSRT (DISCRIMINATIVE CORRELATION FILTER WITH CHANNEL AND SPATIAL RELIABILITY)



1. From left to right: training patch with the bounding box of the object
2. HOG to extract useful information of the image
3. Application of Random Markov Test to generate probabilities
4. Training patch masked using the confidence map

Source: <https://www.arxiv-vanity.com/papers/1611.08461/>

PLAN OF ATTACK

1. Part 1

Biological fundamentals
Single layer perceptron

2. Part 2

Multi-layer perceptron


3. Part 3

Pybrain
Sklearn
TensorFlow
PyTorch


BIOLOGICAL FUNDAMENTALS



BIOLOGICAL FUNDAMENTALS



BIOLOGICAL FUNDAMENTALS



Dendrites


Cell body

Axon

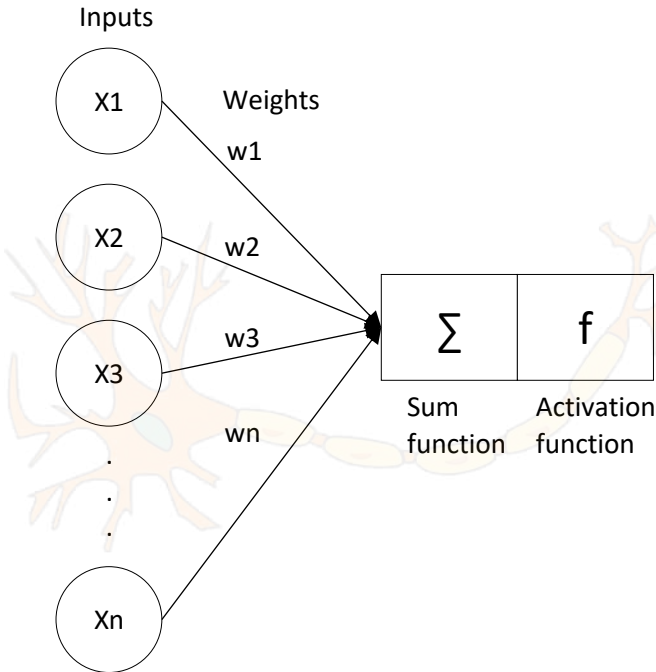
Axon terminals

Synapse

ARTIFICIAL NEURON




ARTIFICIAL NEURON



$$sum = \sum_{i=1}^n x_i * w_i$$

$$x_1 * w_1 + x_2 * w_2 + x_3 * w_3$$

PERCEPTRON



$$sum = \sum_{i=1}^n x_i * w_i$$

$$sum = (35 * 0.8) + (25 * 0.1)$$

$$sum = 28 + 2.5$$


$$sum = 30.5$$

Greater or equal to 1 = 1

Otherwise = 0

“All or nothing” representation

PERCEPTRON



$$sum = \sum_{i=1}^n x_i * w_i$$

$$sum = (35 * -0.8) + (25 * 0.1)$$


$$sum = -28 + 2.5$$

$$sum = -25.5$$

Greater or equal to 1 = 1

Otherwise = 0

STEP FUNCTION







PERCEPTRON

- Positive weight – exciting synapse
- Negative weight – inhibitory synapse
- Weights are the synapses
- Weights amplify or reduce the input signal
- The knowledge of a neural network is the weights

"AND" OPERATOR

X1	X2	Class
0	0	0
0	1	0
1	0	0
1	1	1

"AND" OPERATOR



X1	X2	Class
0	0	0
0	1	0
1	0	0
1	1	1

error = correct - prediction

Class	Prediction	Error
0	0	0
0	0	0
0	0	0
1	0	1


75%

$\text{weight}(n+1) = \text{weight}(n) + (\text{learning_rate} * \text{input} * \text{error})$

"AND" OPERATOR

error = correct - prediction


Class	Prediction	Error
0	0	0
0	0	0
0	0	0
1	0	1



$$\text{weight}(n + 1) = \text{weight}(n) + (\text{learning_rate} * \text{input} * \text{error})$$

$$\text{weight}(n + 1) = 0 + (0.1 * 1 * 1)$$

$$\text{weight}(n + 1) = 0.1$$







$$\text{weight}(n + 1) = \text{weight}(n) + (\text{learning_rate} * \text{input} * \text{error})$$

$$\text{weight}(n + 1) = 0 + (0.1 * 1 * 1)$$

$$\text{weight}(n + 1) = 0.1$$

"AND" OPERATOR



X1	X2	Class
0	0	0
0	1	0
1	0	0
1	1	1

error = correct - prediction


Class	Prediction	Error
0	0	0
0	0	0
0	0	0
1	0	1

75%

"AND" OPERATOR

error = correct - prediction


Class	Prediction	Error
0	0	0
0	0	0
0	0	0
1	0	1



$$\text{weight}(n+1) = \text{weight}(n) + (\text{learning_rate} * \text{input} * \text{error})$$

$$\text{weight}(n+1) = 0.1 + (0.1 * 1 * 1)$$

$$\text{weight}(n+1) = 0.2$$







$$\text{weight}(n+1) = \text{weight}(n) + (\text{learning_rate} * \text{input} * \text{error})$$

$$\text{weight}(n+1) = 0.1 + (0.1 * 1 * 1)$$

$$\text{weight}(n+1) = 0.2$$

"AND" OPERATOR



X1	X2	Class
0	0	0
0	1	0
1	0	0
1	1	1

error = correct - prediction

Class	Prediction	Error
0	0	0
0	0	0
0	0	0
1	1	0

100%

While error $\neq 0$

- For each row

 - Calculate output





 - Calculate error (correct - prediction)

 - If error > 0

 - For each weight

 - Update the weights

"OR" OPERATOR




X1	X2	Class
0	0	0
0	1	1
1	0	1
1	1	1

error = correct - prediction



Class	Prediction	Error
0	0	0
1	1	0
1	1	0
1	1	0

100%

"XOR" OPERATOR

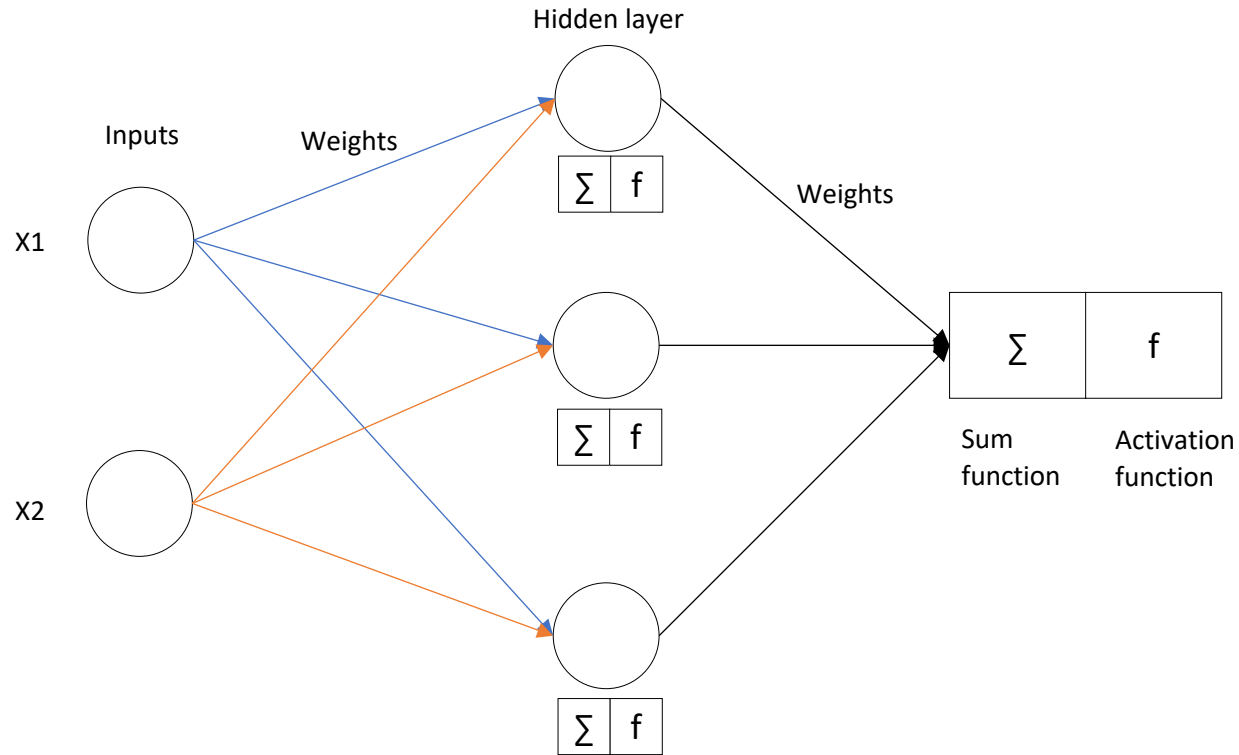


Linearly separable




Non-linearly separable

MULTI-LAYER PERCEPTRON




STEP FUNCTION



SIGMOID FUNCTION


$$y = \frac{1}{1 + e^{-x}}$$



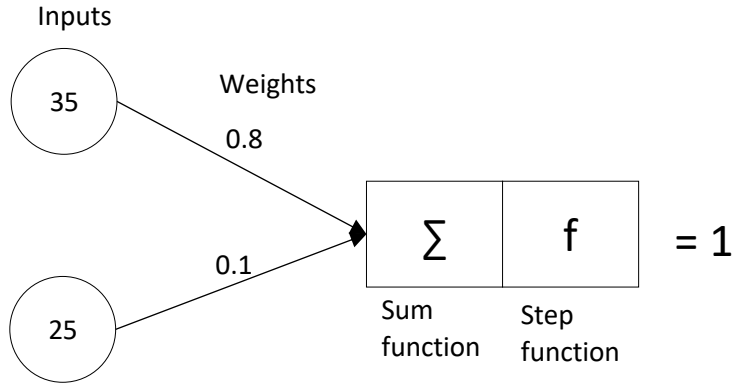
If X is high, the value is approximately 1
If X is small, the value is approximately 0

HYPERBOLIC TANGENT FUNCTION

$$Y = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$



STEP FUNCTION

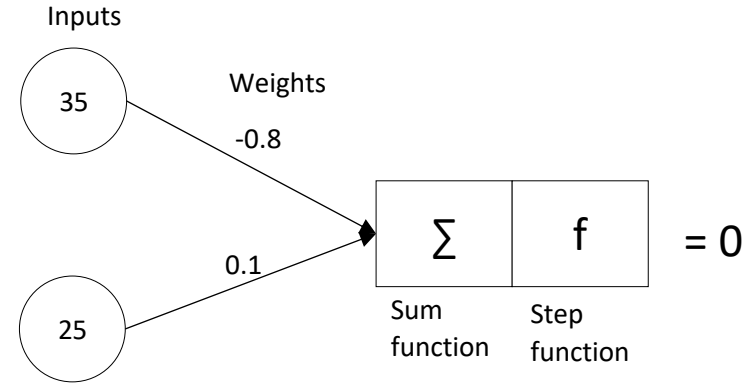


$$sum = \sum_{i=1}^n x_i * w_i$$

$$sum = (35 * 0.8) + (25 * 0.1)$$

$$sum = 28 + 2.5$$

$$sum = 30.5$$



$$sum = \sum_{i=1}^n x_i * w_i$$

$$sum = (35 * -0.8) + (25 * 0.1)$$

$$sum = -28 + 2.5$$

$$sum = -25.5$$

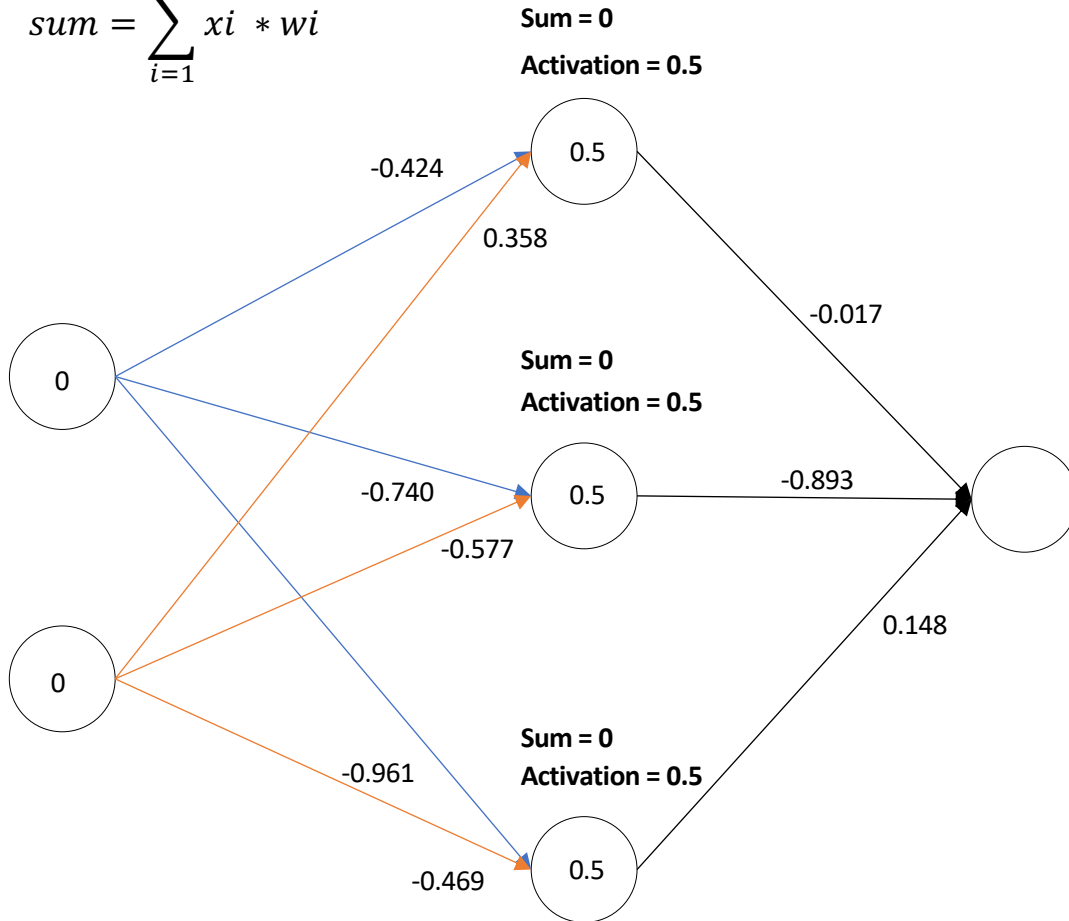
"XOR" OPERATOR

X1	X2	Class
0	0	0
0	1	1
1	0	1
1	1	0

INPUT LAYER TO HIDDEN LAYER



$$sum = \sum_{i=1}^n x_i * w_i$$



$$y = \frac{1}{1 + e^{-x}}$$


X1	X2	Class
0	0	0
0	1	1
1	0	1
1	1	0

$$0 * (-0.424) + 0 * 0.358 = 0$$

$$0 * (-0.740) + 0 * (-0.577) = 0$$

$$0 * (-0.961) + 0 * (-0.469) = 0$$

INPUT LAYER TO HIDDEN LAYER




X1	X2	Class
0	0	0
0	1	1
1	0	1
1	1	0

$$0 * (-0.424) + 1 * 0.358 = 0.358$$

$$0 * (-0.740) + 1 * (-0.577) = -0.577$$

$$0 * (-0.961) + 1 * (-0.469) = -0.469$$

INPUT LAYER TO HIDDEN LAYER




X1	X2	Class
0	0	0
0	1	1
1	0	1
1	1	0

$$1 * (-0.424) + 0 * 0.358 = -0.424$$

$$1 * (-0.740) + 0 * (-0.577) = -0.740$$

$$1 * (-0.961) + 0 * (-0.469) = -0.961$$

INPUT LAYER TO HIDDEN LAYER



X1	X2	Class
0	0	0
0	1	1
1	0	1
1	1	0

$$1 * (-0.424) + 1 * 0.358 = -0.066$$

$$1 * (-0.740) + 1 * (-0.577) = -1.317$$

$$1 * (-0.961) + 1 * (-0.469) = -1.430$$

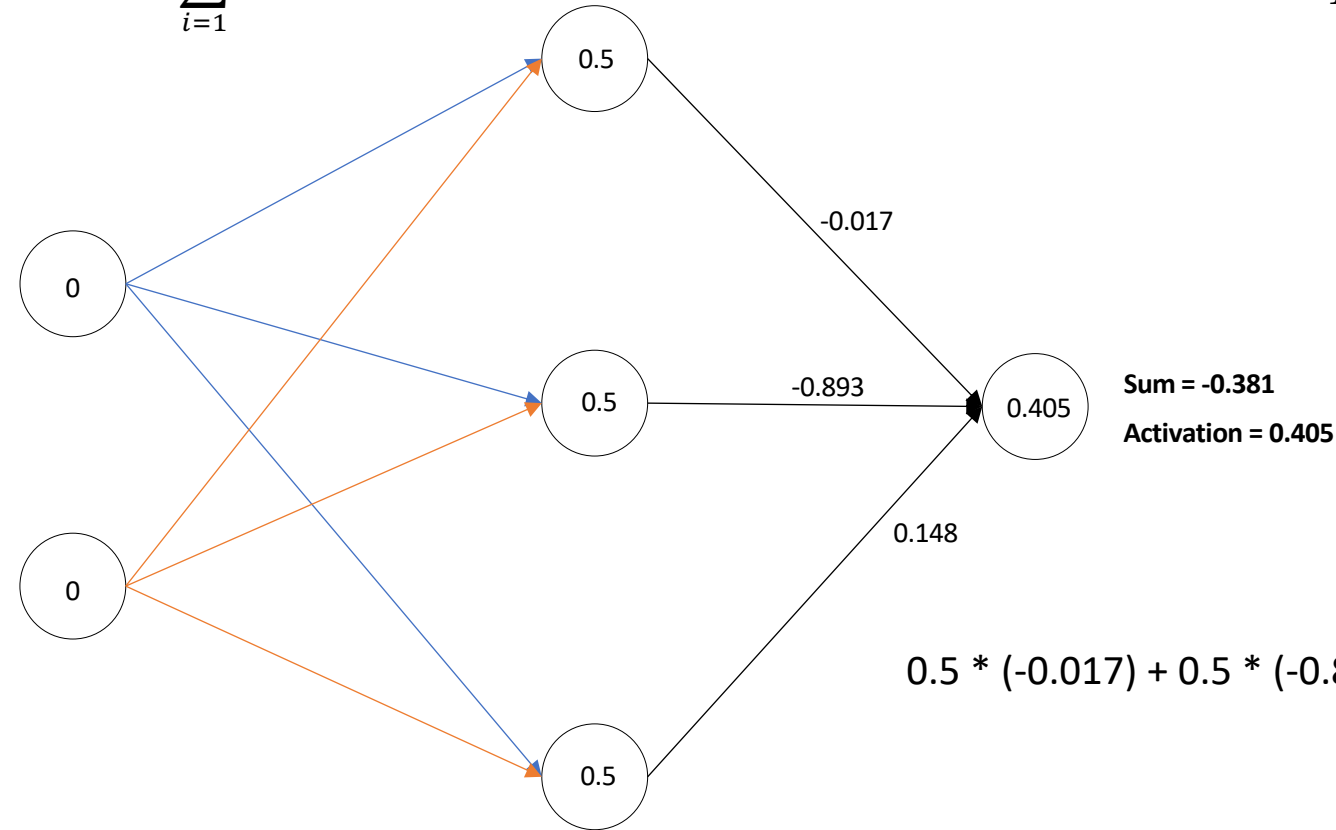
HIDDEN LAYER TO OUTPUT LAYER



$$sum = \sum_{i=1}^n x_i * w_i$$


$$y = \frac{1}{1 + e^{-x}}$$

X1	X2	Class
0	0	0
0	1	1
1	0	1
1	1	0



$$0.5 * (-0.017) + 0.5 * (-0.893) + 0.5 * 0.148 = -0.381$$

HIDDEN LAYER TO OUTPUT LAYER



Sum = -0.274

Activation = 0.431


$$0.589 * (-0.017) + 0.360 * (-0.893) + 0.385 * 0.148 = -0.274$$

X1	X2	Class
0	0	0
0	1	1
1	0	1
1	1	0

HIDDEN LAYER TO OUTPUT LAYER




X1	X2	Class
0	0	0
0	1	1
1	0	1
1	1	0



HIDDEN LAYER TO OUTPUT LAYER



X1	X2	Classe
0	0	0
0	1	1
1	0	1
1	1	0



$$0.483 * (-0.017) + 0.211 * (-0.893) + 0.193 * 0.148 = -0.168$$

"XOR" OPERATOR – ERROR (LOSS FUNCTION)

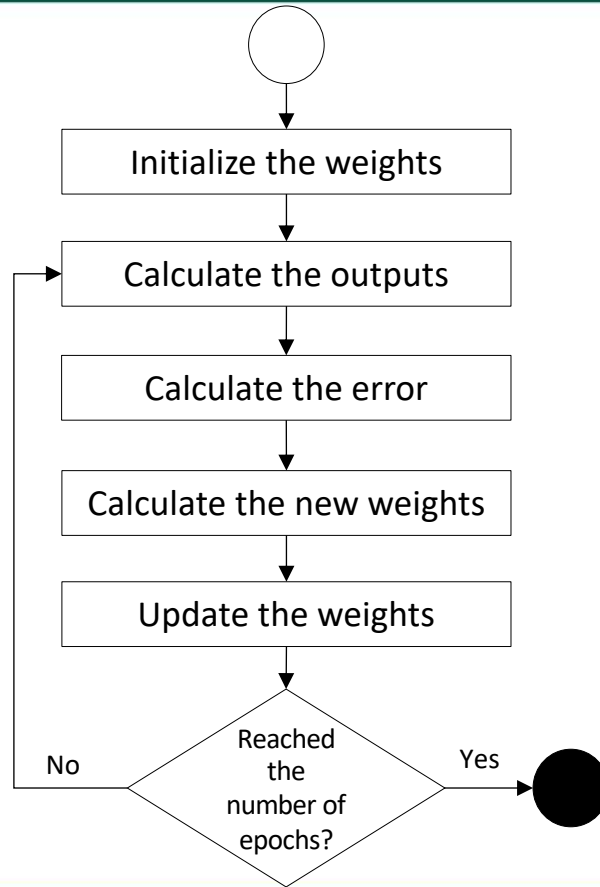
The simplest algorithm

error = correct – prediction

X1	X2	Class	Prediction	Error
0	0	0	0.405	-0.405
0	1	1	0.431	0.569
1	0	1	0.436	0.564
1	1	0	0.458	-0.458

Average = 0.499

ALGORITHM



Cost function (loss function)

Gradient descent

Derivative




Delta

Backpropagation

GRADIENT DESCENT

$$\min C(w_1, w_2 \dots w_n)$$

Calculate the partial derivative to move to the gradient direction



GRADIENT DESCENT (DERIVATIVE)


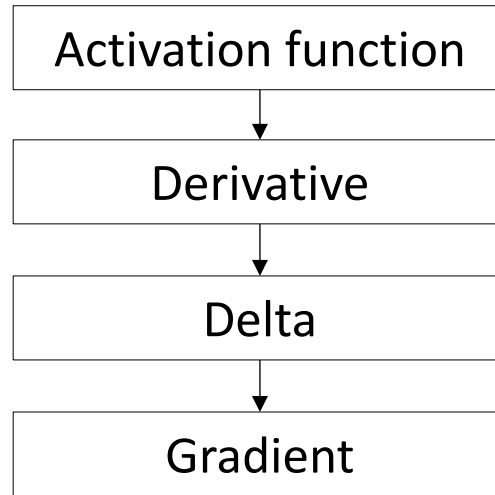
$$y = \frac{1}{1 + e^{-x}}$$



$$d = y * (1 - y)$$

$$d = 0.1 * (1 - 0.1)$$

DELTA PARAMETER




OUTPUT LAYER – DELTA



$$\delta_{output} = error * sigmoid_{derivative}$$

X1	X2	Class
0	0	0
0	1	1
1	0	1
1	1	0



Sum = -0.381

Activation = 0.405

Error = 0 - 0.405 = -0.405

Derivative activation (sigmoid) = 0.241


Delta (output) = -0.405 * 0.241 = -0.097

OUTPUT LAYER – DELTA



$$\delta_{output} = error * sigmoid_{derivative}$$

X1	X2	Class
0	0	0
0	1	1
1	0	1
1	1	0



Sum = -0.274

Activation = 0.431

Error = 1 - 0.431 = 0.569

Derivative activation (sigmoid) = 0.245


Delta (output) = 0.569 * 0.245 = 0.139

OUTPUT LAYER – DELTA



$$\text{delta}_{\text{output}} = \text{error} * \text{sigmoid}_{\text{derivative}}$$

X1	X2	Class
0	0	0
0	1	1
1	0	1
1	1	0



Sum = -0.254

Activation = 0.436

Error = 1 - 0.436 = 0.564

Derivative activation (sigmoid) = 0.246


Delta (output) = 0.564 * 0.246 = 0.138

OUTPUT LAYER – DELTA



$$\text{delta}_{\text{output}} = \text{error} * \text{sigmoid}_{\text{derivative}}$$

X1	X2	Class
0	0	0
0	1	1
1	0	1
1	1	0



Sum = -0.168

Activation = 0.458

Error = 0 - 0.458 = -0.458

Derivative activation (sigmoid) = 0.248

Delta (output) = -0.458 * 0.248 = -0.113

X1	X2	Class
0	0	0
0	1	1
1	0	1
1	1	0



$$0.25 * (-0.893) * (-0.097) = 0.021$$


$$0.25 * 0.148 * (-0.097) = -0.003$$

HIDDEN LAYER – DELTA



$$\delta_{hidden} = \text{sigmoid}_{derivative} * \text{weight} * \delta_{output}$$

X1	X2	Class
0	0	0
0	1	1
1	0	1
1	1	0



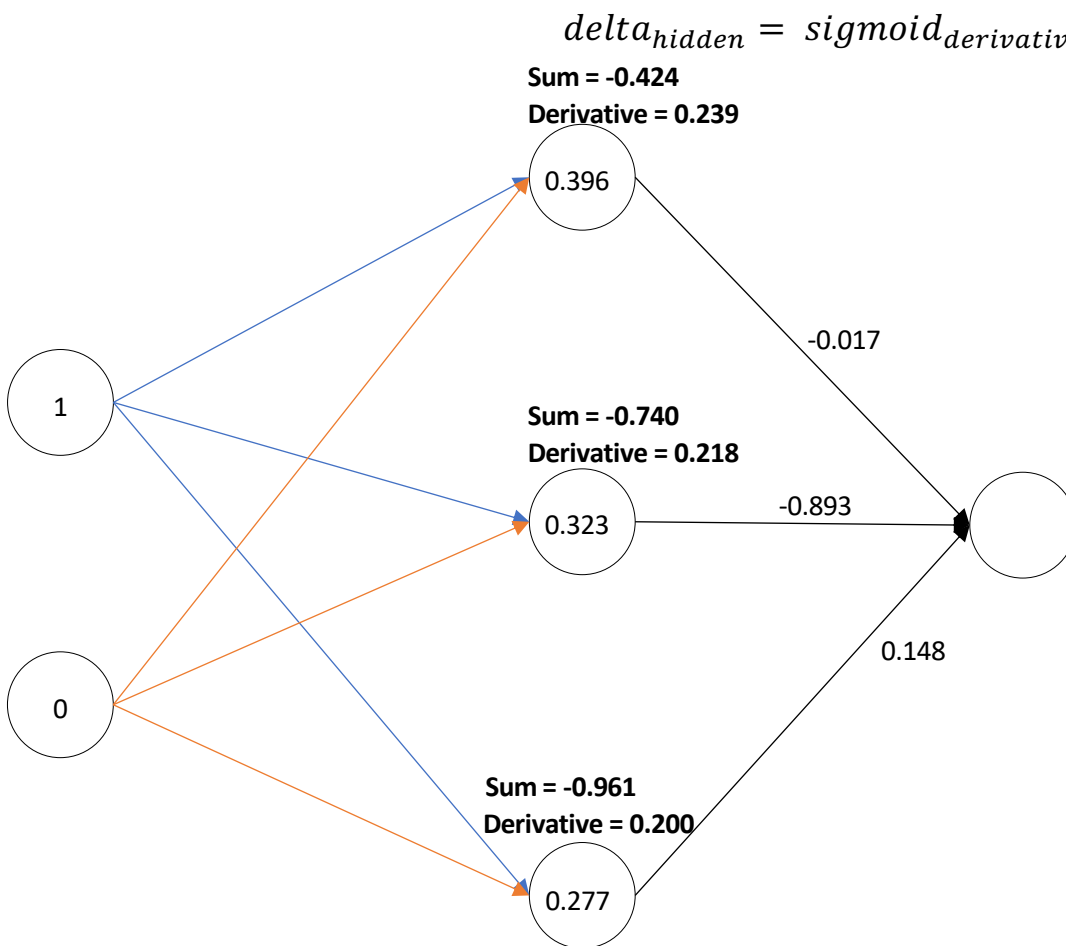
Delta (output) = 0.139

$$0.242 * (-0.017) * 0.139 = -0.000$$

$$0.230 * (-0.893) * 0.139 = -0.028$$

$$0.236 * 0.148 * 0.139 = 0.004$$

HIDDEN LAYER – DELTA



X1	X2	Class
0	0	0
0	1	1
1	0	1
1	1	0

Delta (output) = 0.138

$$0.239 * (-0.017) * 0.138 = -0.000$$

$$0.218 * (-0.893) * 0.138 = -0.026$$


$$0.200 * 0.148 * 0.138 = 0.004$$

HIDDEN LAYER – DELTA



$$\delta_{hidden} = \text{sigmoid}_{derivative} * \text{weight} * \delta_{output}$$

X1	X2	Class
0	0	0
0	1	1
1	0	1
1	1	0




Delta (output) = -0.113

$$0.249 * (-0.017) * (-0.113) = 0.000$$

$$0.166 * (-0.893) * (-0.113) = 0.016$$

$$0.155 * 0.148 * (-0.113) = -0.002$$


WEIGHT UPDATE



$$weight_{n+1} = weight_n + (input * delta * learning_rate)$$

LEARNING RATE



- Defines how fast the algorithm will learn
- High: the convergence is fast but may lose the global minimum
- Low: the convergence will be slower but more likely to reach the global minimum





WEIGHT UPDATE – OUTPUT LAYER TO HIDDEN LAYER



$$weight_{n+1} = weight_n + (input * delta * learning_rate)$$





$$0.5 * (-0.097) + 0.588 * 0.139 + 0.395 * 0.138 + 0.483 * (-0.113) = 0.033$$





WEIGHT UPDATE – OUTPUT LAYER TO HIDDEN LAYER



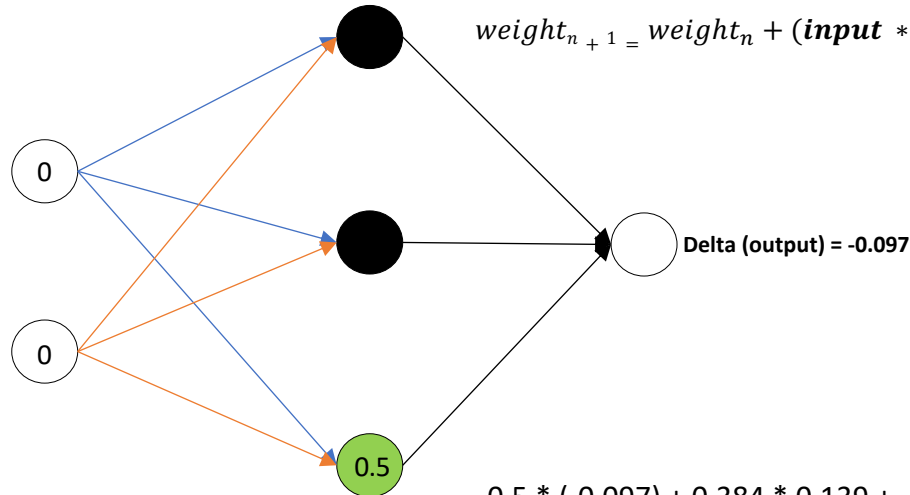
$$weight_{n+1} = weight_n + (input * delta * learning_rate)$$






$$0.5 * (-0.097) + 0.359 * 0.139 + 0.323 * 0.138 + 0.211 * (-0.113) = \mathbf{0.022}$$



WEIGHT UPDATE – OUTPUT LAYER TO HIDDEN LAYER



$$0.5 * (-0.097) + 0.384 * 0.139 + 0.276 * 0.138 + 0.193 * (-0.113) = \mathbf{0.021}$$



WEIGHT UPDATE – OUTPUT LAYER TO HIDDEN LAYER

Learning rate = 0.3

Input x delta

0.033

0.022


0.021

$$weight_{n+1} = weight_n + (input * delta * learning_rate)$$



$$-0.017 + 0.033 * 0.3 = \mathbf{-0.007}$$

$$-0.893 + 0.022 * 0.3 = \mathbf{-0.886}$$



$$0.148 + 0.021 * 0.3 = \mathbf{0.154}$$



WEIGHT UPDATE – HIDDEN LAYER TO INPUT LAYER



$$weight_{n+1} = weight_n + (input * delta * learning_rate)$$







$$0 * 0.000 + 0 * (-0.001) + 1 * (0.001) + 1 * -0.000 = -0.000$$

$$0 * 0.000 + 1 * (-0.001) + 0 * (0.001) + 1 * -0.000 = -0.000$$

WEIGHT UPDATE – HIDDEN LAYER TO INPUT LAYER



$$weight_{n+1} = weight_n + (input * delta * learning_rate)$$







$$0 * 0.021 + 0 * (-0.028) + 1 * (-0.026) + 1 * 0.016 = -0.009$$

$$0 * 0.021 + 1 * (-0.028) + 0 * (-0.026) + 1 * 0.016 = -0.012$$

WEIGHT UPDATE – HIDDEN LAYER TO INPUT LAYER



$$weight_{n+1} = weight_n + (input * delta * learning_rate)$$



$$0 * (-0.003) + 0 * 0.004 + 1 * 0.004 + 1 * (-0.002) = 0.002$$

$$0 * (-0.003) + 1 * 0.004 + 0 * 0.004 + 1 * (-0.002) = 0.002$$

WEIGHT UPDATE – HIDDEN LAYER TO INPUT LAYER



Learning rate = 0.3

Input x delta

-0.000 -0.009 0.002

-0.000 -0.012 0.002

$$weight_{n+1} = weight_n + (input * delta * learning_rate)$$

$$-0.424 + (-0.000) * 0.3 = \mathbf{-0.424}$$

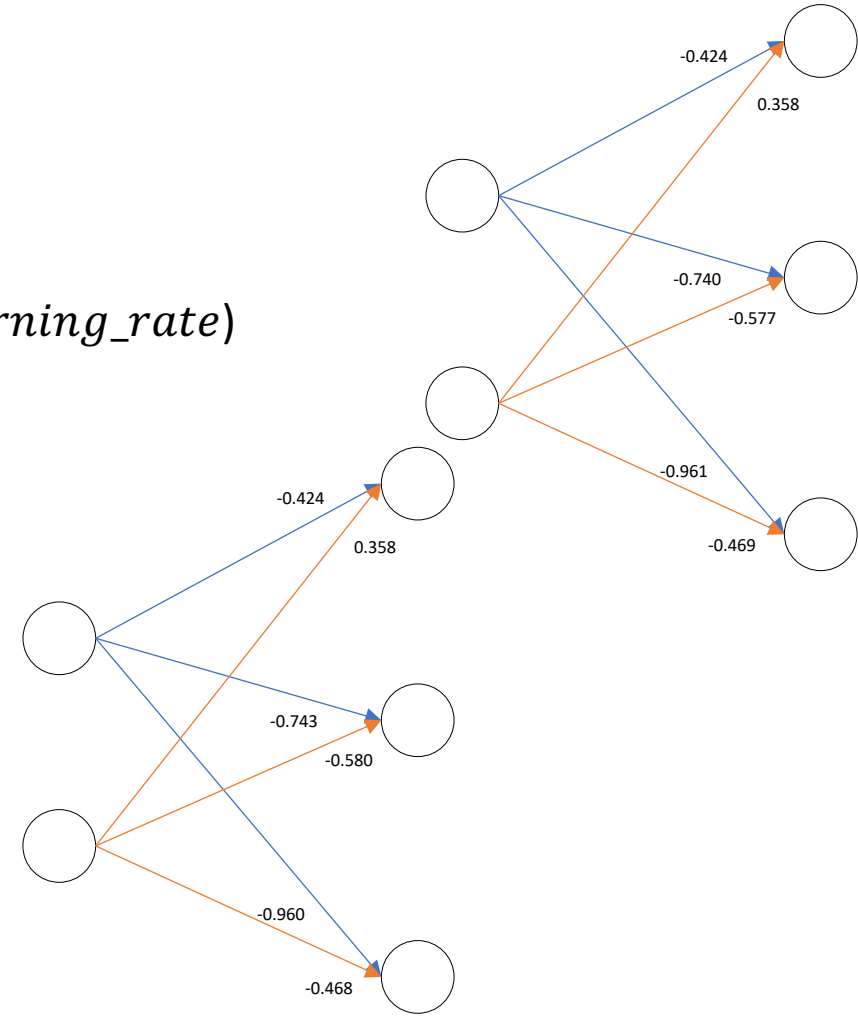
$$0.358 + (-0.000) * 0.3 = \mathbf{0.358}$$

$$-0.740 + (-0.009) * 0.3 = \mathbf{-0.743}$$

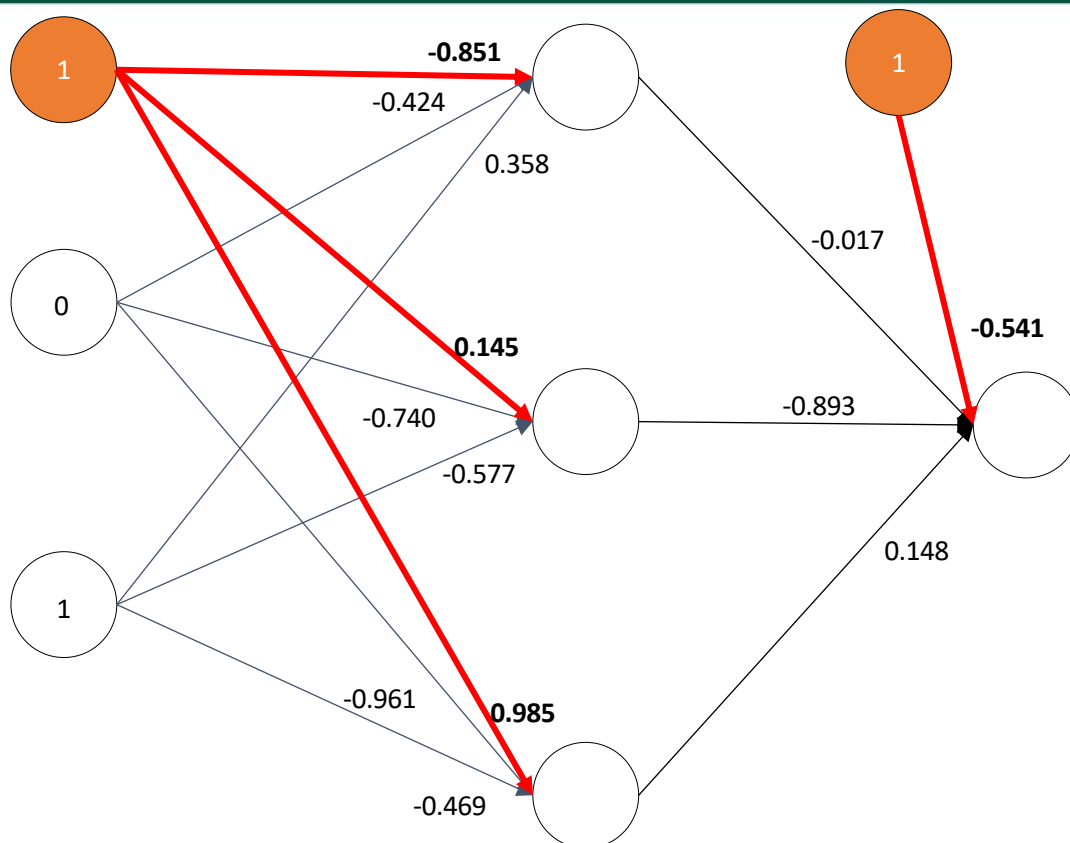
$$-0.577 + (-0.012) * 0.3 = \mathbf{-0.580}$$

$$-0.961 + 0.002 * 0.3 = \mathbf{-0.960}$$

$$-0.469 + 0.002 * 0.3 = \mathbf{-0.468}$$



BIAS



$$output = \sum(inputs * weights) + bias$$

ERROR (LOSS FUNCTION)

The simplest algorithm

error = correct – prediction

X1	X2	Class	Prediction	Error
0	0	0	0.405	-0.405
0	1	1	0.431	0.569
1	0	1	0.436	0.564
1	1	0	0.458	-0.458

Average = 0.499

MEAN SQUARED ERROR (MSE) AND ROOT MEAN SQUARED ERROR (RMSE)

$$MSE = \frac{1}{N} \sum_{i=1}^N (f_i - y_i)^2$$

$$RMSE = \sqrt{\frac{1}{N} \sum_{i=1}^N (f_i - o_i)^2}$$

X1	X2	Class	Prediction	Error
0	0	0	0.405	$(0 - 0.405)^2 = 0.164$
0	1	1	0.431	$(1 - 0.431)^2 = 0.322$
1	0	1	0.436	$(1 - 0.436)^2 = 0.316$
1	1	0	0.458	$(0 - 0.458)^2 = 0.209$

10 (expected) – 5 (prediction) = 5 ($5^2 = 25$)

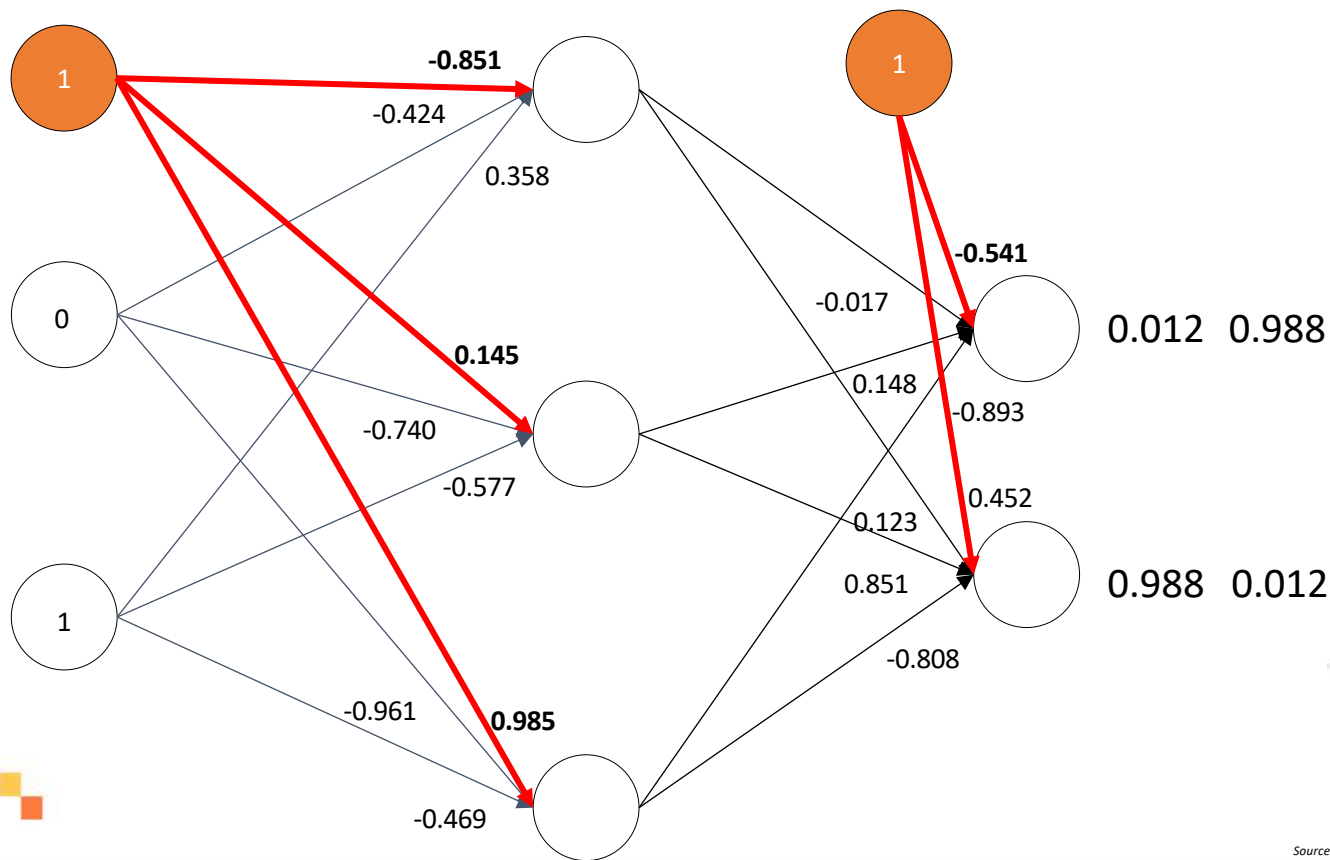
10 (expected) – 8 (prediction) = 2 ($2^2 = 4$)

Sum = 1.011

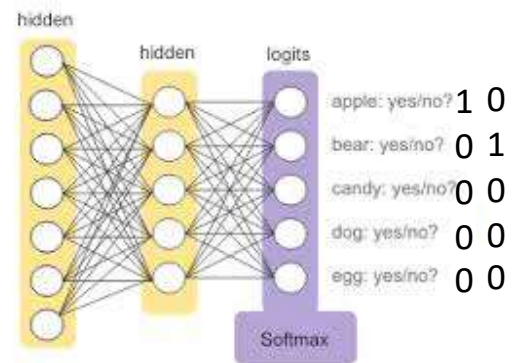
MSE = 1.011 / 4 = 0.252

RMSE = 0.501

MULTIPLE OUTPUTS



X1	X2	Class
0	0	0
0	1	1
1	0	1
1	1	0



Source: <https://developers.google.com/machine-learning/crash-course/multi-class-neural-networks/softmax>


HIDDEN LAYERS




$$\text{Neurons} = \frac{\text{Inputs} + \text{Outputs}}{2}$$

Inputs


Output



$$\text{Neurons} = \frac{2 + 1}{2} = 1.5$$



$$\text{Neurons} = \frac{7 + 1}{2} = 4$$



$$\text{Neurons} = \frac{8 + 2}{2} = 5$$

HIDDEN LAYERS

- Linearly separable problems do not require hidden layers
- In general, two layers work well
- Deep learning research shows that more layers are essential for more complex problems

HIDDEN LAYERS






Credit history	Debts	Properties	Anual income	Risk
Bad	High	No	< 15.000	High
Unknown	High	No	>= 15.000 a <= 35.000	High
Unknown	Low	No	>= 15.000 a <= 35.000	Moderate
Unknown	Low	No	> 35.000	High
Unknown	Low	No	> 35.000	Low
Unknown	Low	Yes	> 35.000	Low
Bad	Low	No	< 15.000	High
Bad	Low	Yes	> 35.000	Moderate
Good	Low	No	> 35.000	Low
Good	High	Yes	> 35.000	Low
Good	High	No	< 15.000	High
Good	High	No	>= 15.000 a <= 35.000	Moderate
Good	High	No	> 35.0000	Low
Bad	High	No	>= 15.000 a <= 35.000	High

HIDDEN LAYERS




$$\text{Neurons} = \frac{4 + 3}{2} = 3,5$$



The higher the activation value,
the more impact the neuron has

OUTPUT LAYER WITH CATEGORICAL DATA



error = correct – prediction

expected output = 1 0 0

prediction = 0.95 0.02 0.03

error = (1 – 0.95) + (0 – 0.02) + (0 – 0.03)

error = 0.05 + 0.02 + 0.03 = 0.08

Credit history	Debts	Properties	Annual income	Risk
3	1	1	1	100
2	1	1	2	100
2	2	1	2	010
2	2	1	3	100
2	2	1	3	001
2	2	2	3	001
3	2	1	1	100
3	2	2	3	010
1	2	1	3	001
1	1	2	3	001
1	1	1	1	100
1	1	1	2	010
1	1	1	3	001
3	1	1	2	100

STOCHASTIC GRADIENT DESCENT



Credit history	Debts	Properties	Anual income	Risk
3	1	1	1	100
2	1	1	2	100
2	2	1	2	010
2	2	1	3	100
2	2	1	3	001
2	2	2	3	001
3	2	1	1	100
3	2	2	3	010
1	2	1	3	001
1	1	2	3	001
1	1	1	1	100
1	1	1	2	010
1	1	1	3	001
3	1	1	2	100


Batch gradient descent

Calculate the error for all instances and then update the weights


Stochastic gradient descent

Calculate the error for each instance and then update the weights


Credit history	Debts	Properties	Anual income	Risk
3	1	1	1	100
2	1	1	2	100
2	2	1	2	010
2	2	1	3	100
2	2	1	3	001
2	2	2	3	001
3	2	1	1	100
3	2	2	3	010
1	2	1	3	001
1	1	2	3	001
1	1	1	1	100
1	1	1	2	010
1	1	1	3	001
3	1	1	2	100



CONVEX AND NON CONVEX




Convex



Non convex

- Stochastic gradient descent
 - Prevent local minimums (non convex)
 - Faster
- Mini batch gradient descent
 - Select a pre-defined number of instances in order to calculate the error and update the weights


- 90's: SVM (Support Vector Machines)
- From 2006, several algorithms were created for training neural networks
- Two or more hidden layers



- Convolutional neural networks
- Recurrent neural networks
- Autoencoders
- GANs (Generative adversarial networks)


SOFTMAX

$$Y = \frac{e(x)}{\sum e(x)}$$

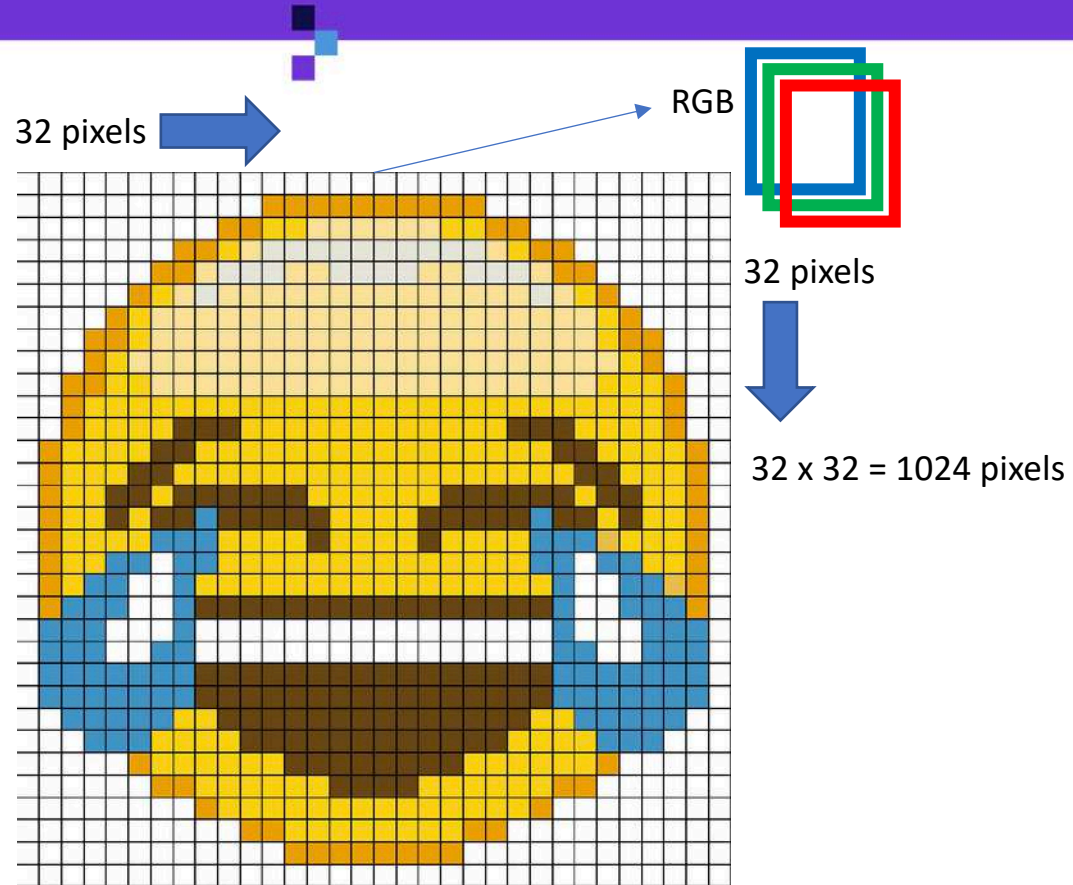


PLAN OF ATTACK – NEURAL NETWORKS FOR IMAGE CLASSIFICATION

1. Intuition about neural networks
2. Neural network using the pixels of the images
3. Feature extractor with OpenCV
4. Neural network using feature extraction




PIXELS





PIXELS

1	2	3	4	5	6	7	8	9	0
1	2	3	4	5	6	7	8	9	0
1	2	3	4	5	6	7	8	9	0
1	2	3	4	5	6	7	8	9	0
1	2	3	4	5	6	7	8	9	0
1	2	3	4	5	6	7	8	9	0
1	2	3	4	5	6	7	8	9	0
1	2	3	4	5	6	7	8	9	0
1	2	3	4	5	6	7	8	9	0
1	2	3	4	5	6	7	8	9	0



0,1,1,1,0,0,1,0,0,0,0,1,1,1,0,0,1,0,1,0,0,1,1,1,0,6



0,1,1,1,0,0,1,0,1,0,0,1,1,1,0,0,0,0,1,0,0,1,1,1,0,9

FEATURE EXTRACTION

Brown (mouth)

Blue (pants)

Gray (shoes)




Orange (T-Shirt)

Blue (shorts)

Blue (sneakers)




0.0,0.0,0.087449393,9.861538462,8.660728745,2.257489879,0
0.0,0.0,0.012995452,4.980506823,0.0,0.0,0
0.0,0.0,0.014219662,9.668452534,1.034365711,1.034365711,0
0.0,0.0,0.013168984,10.734478157,5.981352718,2.126352016,0
0.979401078,16.958948664,2.349976559,2.48769339,0.0,0.0,1
0.058665773,15.566543748,0.306738183,0.0,0.0,0.0,1
1.829063147,14.074792961,2.135093168,0.0,0.0,0.0,1
2.804717308,0.152271939,0.026014568,0.0,0.0,0.0,1
2.510348847,0.0,0.037677839,0.0,0.0,0.0,1






PLAN OF ATTACK – CONVOLUTIONAL NEURAL NETWORKS FOR IMAGE CLASSIFICATION

1. Intuition about convolutional neural networks
2. Types of neural networks for image classification
3. Convolutional neural network for image classification



PIXELS

$32 \times 32 = 1.024 \times 3 = 3.072$ inputs



- It does not use all pixels
- It applies a dense neural network, but at the beginning it transforms the data
- What are the most important features?

CONVOLUTIONAL NEURAL NETWORK STEPS

1. Convolution operation
2. Pooling
3. Flattening
4. Dense neural network

STEP 1: CONVOLUTION OPERATION

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	1	1
0	1	0	1	1	0	0
0	1	0	1	1	0	1
0	1	0	0	0	1	1

Image

x

1	0	0
1	0	1
0	1	1

Feature detector

=

0				

Feature map

$$0 * 1 + 0 * 0 + 0 * 0 + 0 * 1 + 1 * 0 + 0 * 1 + 0 * 0 + 0 * 1 + 0 * 1 = 0$$

STEP 1: CONVOLUTION OPERATION

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	1	1
0	1	0	1	1	0	0
0	1	0	1	1	0	1
0	1	0	0	0	1	1

Image

x

1	0	0
1	0	1
0	1	1

Feature detector

=

0	1			

Feature map

$$0 * 1 + 0 * 0 + 0 * 0 + 1 * 1 + 0 * 0 + 0 * 1 + 0 * 0 + 0 * 1 + 0 * 1 = 1$$

STEP 1: CONVOLUTION OPERATION



0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	1	1
0	1	0	1	1	0	0
0	1	0	1	1	0	1
0	1	0	0	0	1	1

Image

x

1	0	0
1	0	1
0	1	1

Feature detector

=

0	1	0		

Feature map

$$0 * 1 + 0 * 0 + 0 * 0 + 0 * 1 + 0 * 0 + 0 * 1 + 0 * 0 + 0 * 1 + 0 * 1 = 0$$



STEP 1: CONVOLUTION OPERATION

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	1	1
0	1	0	1	1	0	0
0	1	0	1	1	0	1
0	1	0	0	0	1	1

Image

x

1	0	0
1	0	1
0	1	1

Feature detector

=

0	1	0	1	

Feature map

$$0 * 1 + 0 * 0 + 0 * 0 + 0 * 1 + 0 * 0 + 1 * 1 + 0 * 0 + 0 * 1 + 0 * 1 = 1$$

STEP 1: CONVOLUTION OPERATION

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	1	1
0	1	0	1	1	0	0
0	1	0	1	1	0	1
0	1	0	0	0	1	1

Image

x

1	0	0
1	0	1
0	1	1

Feature detector

=

0	1	0	1	0
0	2	1	1	2
1	2	2	3	1
1	3	3	3	2
1	3	1	3	5

Feature map

$$1 * 1 + 0 * 0 + 0 * 0 + 1 * 1 + 0 * 0 + 1 * 1 + 0 * 0 + 1 * 1 + 1 * 1 = 5$$

STEP 1: CONVOLUTION OPERATION – RELU

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	1	1
0	1	0	1	1	0	0
0	1	0	1	1	0	1
0	1	0	0	0	1	1

Image

x



1	0	0
1	0	1
0	1	1

Feature detector

=

0	1	0	1	0
0	2	1	1	2
1	2	2	3	1
1	3	3	3	2
1	3	1	3	5

Feature map



STEP 1: CONVOLUTIONAL LAYER

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	1	1
0	1	0	1	1	0	0
0	1	0	1	1	0	1
0	1	0	0	0	1	1

Image




0	1	0	1	0
0	2	1	1	2
1	2	2	3	1
1	3	3	3	2
1	3	1	3	5

Feature maps

The network will decide which feature detector to use

The convolutional layer is the set of feature maps

STEP 2: POOLING



STEP 2: POOLING

0	1	0	1	0
0	2	1	1	2
1	2	2	3	1
1	3	3	3	2
1	3	1	3	5

Feature map



2		

STEP 2: POOLING

0	1	0	1	0
0	2	1	1	2
1	2	2	3	1
1	3	3	3	2
1	3	1	3	5

Feature map



2	1	

STEP 2: POOLING

0	1	0	1	0
0	2	1	1	2
1	2	2	3	1
1	3	3	3	2
1	3	1	3	5

Feature map



2	1	2

STEP 2: POOLING

0	1	0	1	0
0	2	1	1	2
1	2	2	3	1
1	3	3	3	2
1	3	1	3	5

Feature map



2	1	2
3		

STEP 2: POOLING

0	1	0	1	0
0	2	1	1	2
1	2	2	3	1
1	3	3	3	2
1	3	1	3	5

Feature map



2	1	2
3	3	

STEP 2: POOLING

0	1	0	1	0
0	2	1	1	2
1	2	2	3	1
1	3	3	3	2
1	3	1	3	5

Feature map



2	1	2
3	3	2

STEP 2: POOLING

0	1	0	1	0
0	2	1	1	2
1	2	2	3	1
1	3	3	3	2
1	3	1	3	5

Feature map



2	1	2
3	3	2
3		

STEP 2: POOLING

0	1	0	1	0
0	2	1	1	2
1	2	2	3	1
1	3	3	3	2
1	3	1	3	5

Feature map



2	1	2
3	3	2
3	3	

STEP 2: POOLING


0	1	0	1	0
0	2	1	1	2
1	2	2	3	1
1	3	3	3	2
1	3	1	3	5

Feature map



2	1	2
3	3	2
3	3	5

CONVOLUTIONAL NEURAL NETWORK – POOLING




STEP 3: FLATTENING

2	1	2
3	3	2
3	3	5

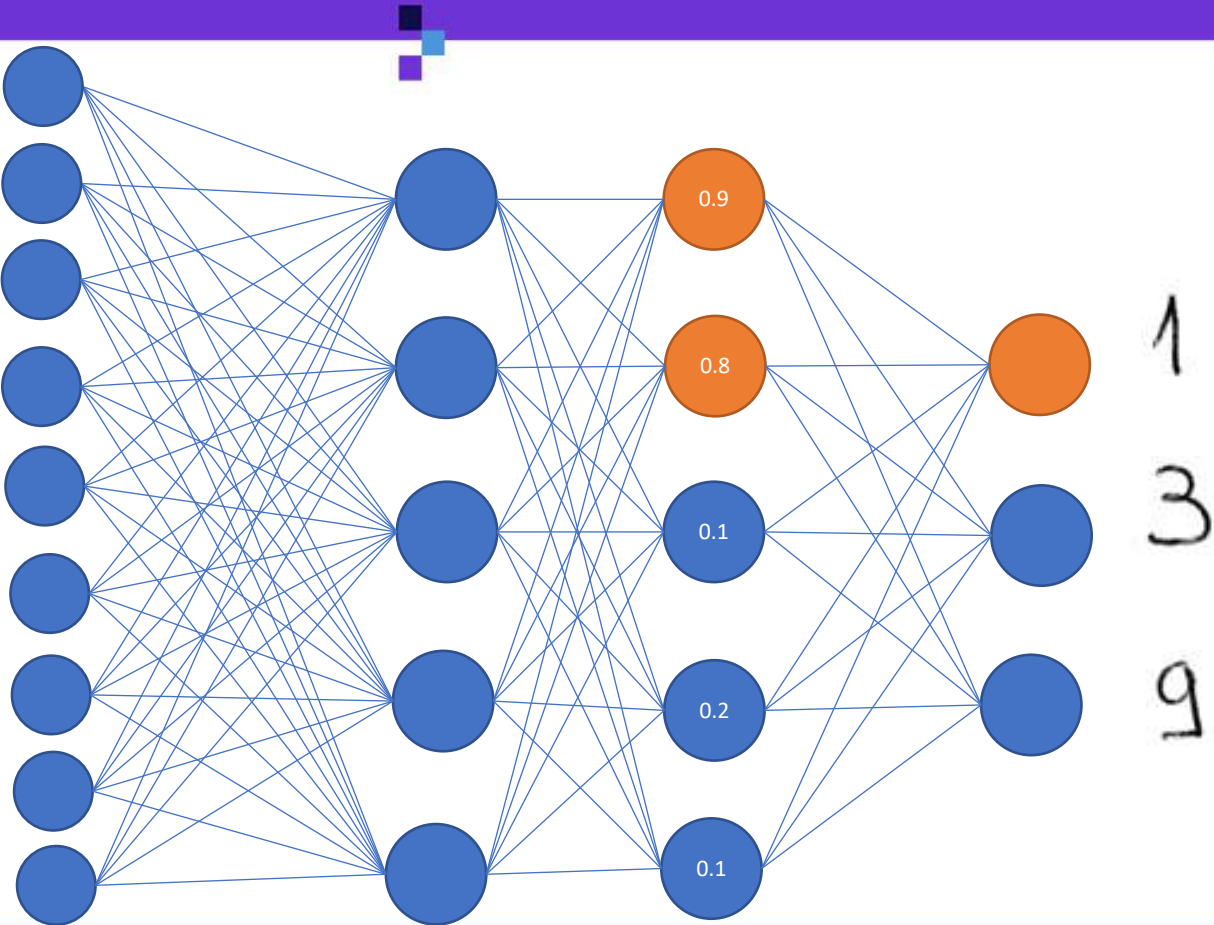
Pooled feature map



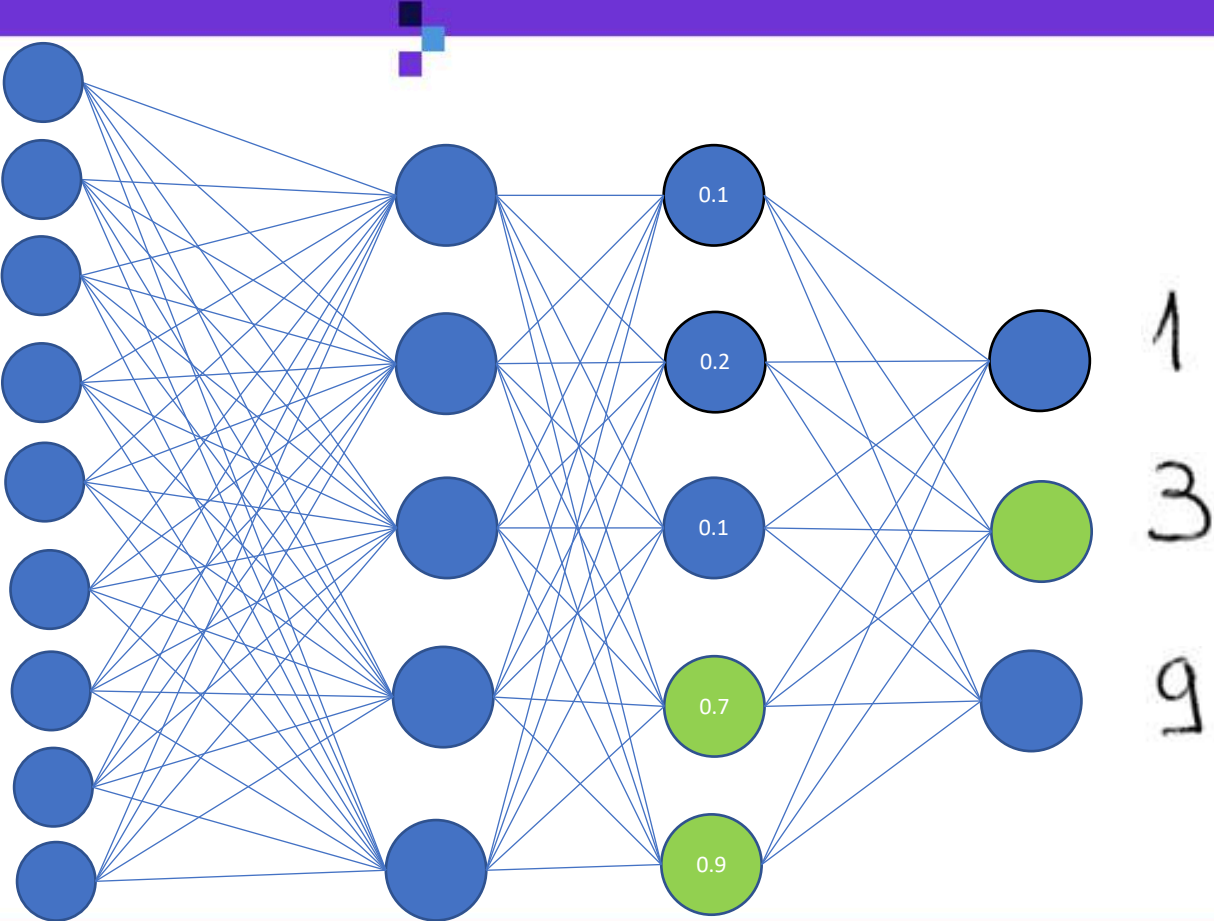
2
1
2
3
3
2
3
3
5



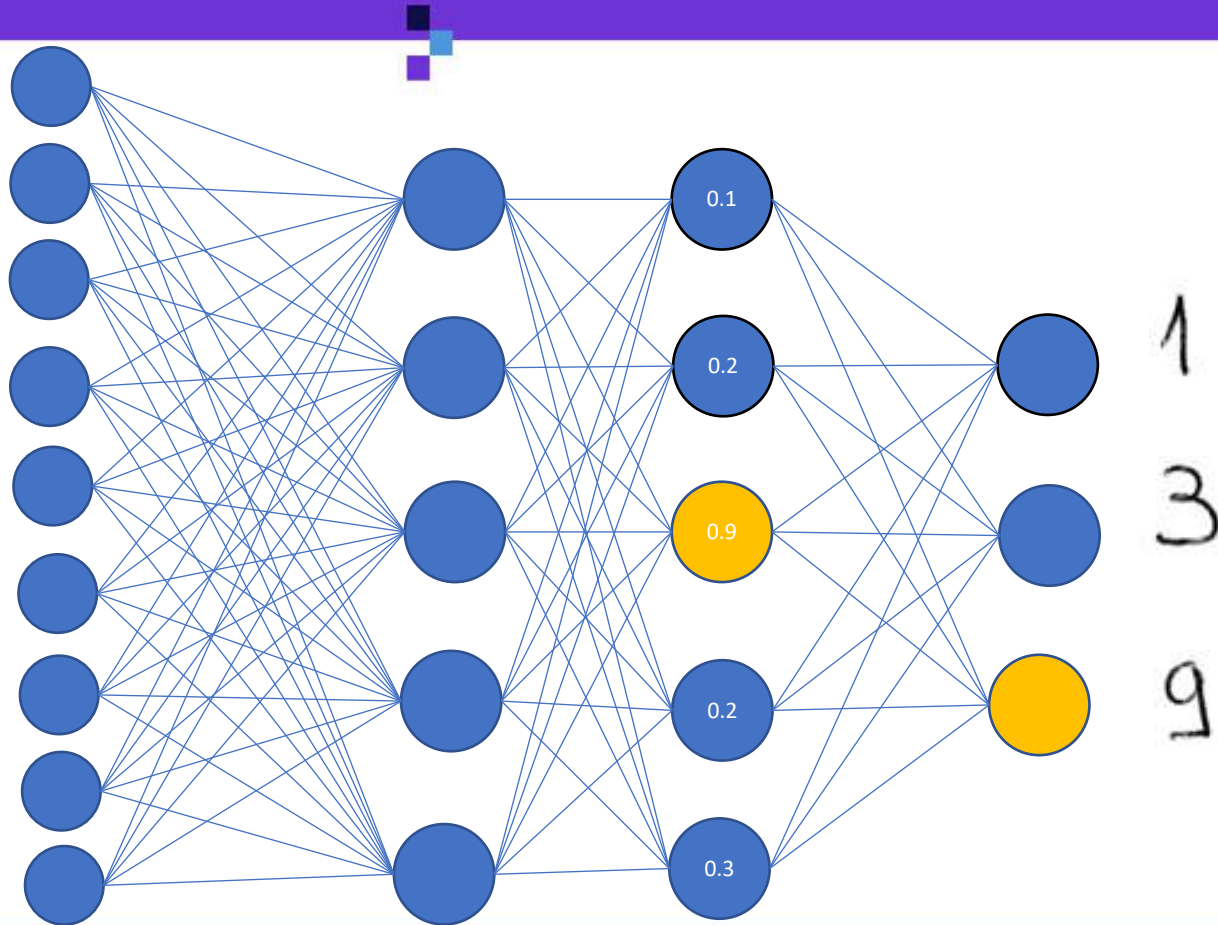
STEP 3: DENSE NEURAL NETWORK



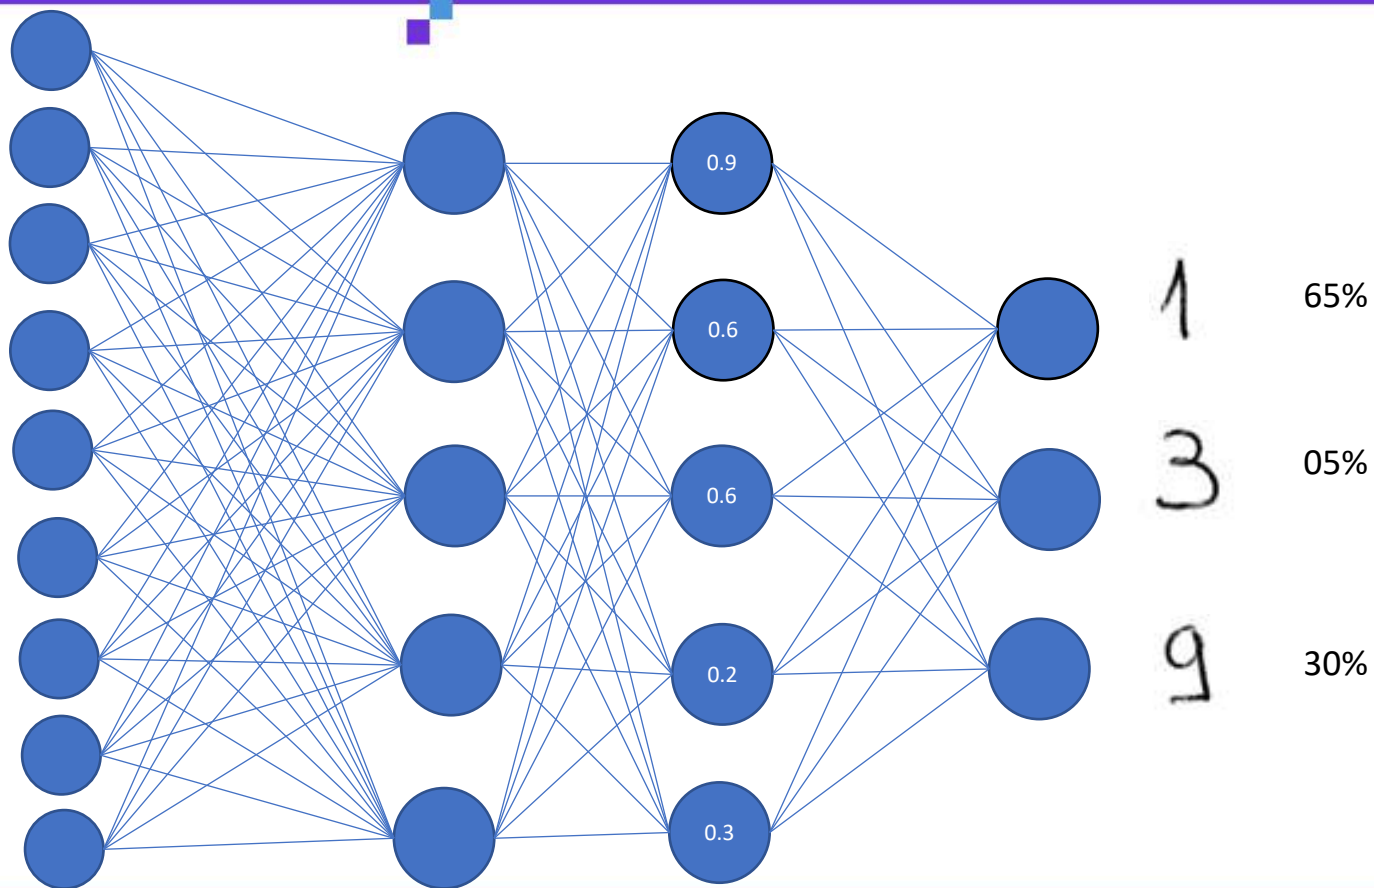
STEP 3: DENSE NEURAL NETWORK





STEP 3: DENSE NEURAL NETWORK



STEP 3: DENSE NEURAL NETWORK



CONVOLUTIONAL NEURAL NETWORK




Training using gradient descent

In addition to adjusting the weights, the feature detector is also changed



PLAN OF ATTACK – TRANSFER LEARNING AND FINE TUNING FOR IMAGE CLASSIFICATION

1. Intuition about transfer learning
2. Implementation
3. Intuition about fine tuning
4. Implementation




TRANSFER LEARNING




KNOWLEDGE TRANSFER



TRANSFER LEARNING



FINE TUNING



PLAN OF ATTACK – NEURAL NETWORKS FOR CLASSIFICATION OF EMOTIONS




1. Implementation of convolutional neural networks
2. Detecting emotions in images
3. Detecting emotions in videos




Source: EPFL

PLAN OF ATTACK – AUTOENCODERS


1. Intuition about autoencoders
2. Implementation of linear autoencoders
3. Implementation of convolutional autoencoders




AUTOENCODERS



LINEAR AUTOENCODERS




LINEAR AUTOENCODERS



Source: <https://skymind.ai/wiki/deep-autoencoder>



CONVOLUTIONAL AUTOENCODERS



Source: <https://www.semanticscholar.org/paper/Deep-Clustering-with-Convolutional-Autoencoders-Guo-Liu/b4c8c77fe8ac7aa07a9e41827955d38a24f6137f>



AUTOENCODERS – APPLICATIONS

1. Noise removal
2. Image compression
3. Fraud detection
4. Dimensionality reduction (similar to PCA)

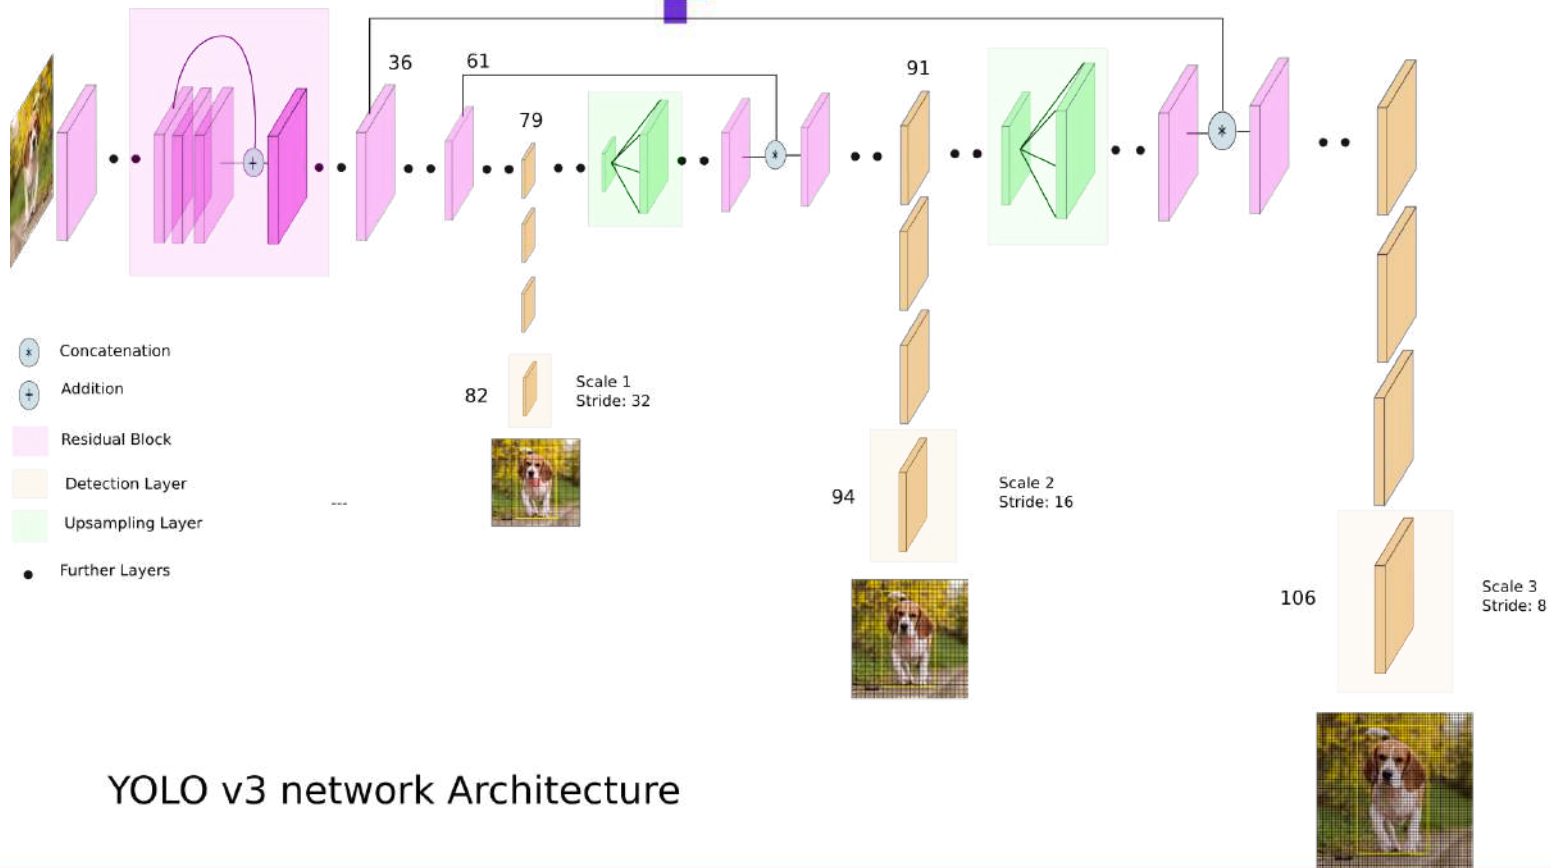


PLAN OF ATTACK – YOLO (YOU ONLY LOOK ONCE)

1. Intuition about YOLO
2. Object detection in images (Darknet)
3. Object detection in videos




YOLO (YOU ONLY LOOK ONCE)



YOLO v3 network Architecture

PLAN OF ATTACK – RECOGNITION OF GESTURES AND ACTIONS

1. Intuition about recognition of gestures and actions
2. Implementation using images and videos




MPII MODEL



mpi

max planck institut
informatik


<http://human-pose.mpi-inf.mpg.de/>



MPII – BODY POINTS



0	Head
1	Neck
2	Rigth shoulder
3	Rigth elbow
4	Rigth wrist
5	Left shoulder
6	Left elbow
7	Left wrist
8	Rigth hip

9	Rigth knee
10	Rigth ankle
11	Left hip
12	Left knee
13	Left ankle
14	Chest
15	Background





PLAN OF ATTACK – DEEP DREAM




1. Intuition about Deep Dream
2. Implementation



DEEP DREAM



General features



Classification specific layers


PLAN OF ATTACK – STYLE TRANSFER

1. Intuition about Style Transfer
2. Implementation



Source of images: https://www.cv-foundation.org/openaccess/content_cvpr_2016/papers/Gatys_Image_Style_Transfer_CVPR_2016_paper.pdf

STYLE TRANSFER



STYLE TRANSFER



PLAN OF ATTACK – GANs (GENERATIVE ADVERSARIAL NETWORKS)

1. Intuition about GANs
2. Implementation

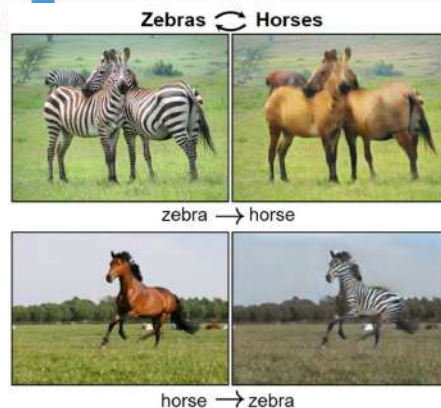




Figure 5: 1024 × 1024 images generated using the CELEBA-HQ dataset. See Appendix F for a larger set of results, and the accompanying video for latent space interpolations.




Source of images: <https://jonathan-hui.medium.com/gan-some-cool-applications-of-gans-4c9ecca35900>




GANs (GENERATIVE ADVERSARIAL NETWORKS)

GENERATOR



DISCRIMINATOR




GANs (GENERATIVE ADVERSARIAL NETWORKS)



BANK



REAL MONEY




DISCRIMINATOR

GENERATOR





FAKE MONEY



GANs – GENERATOR

1 – Real images
0 – Fake images

GENERATOR



GANs - DISCRIMINATOR




FAKE MONEY



REAL MONEY

DISCRIMINATOR

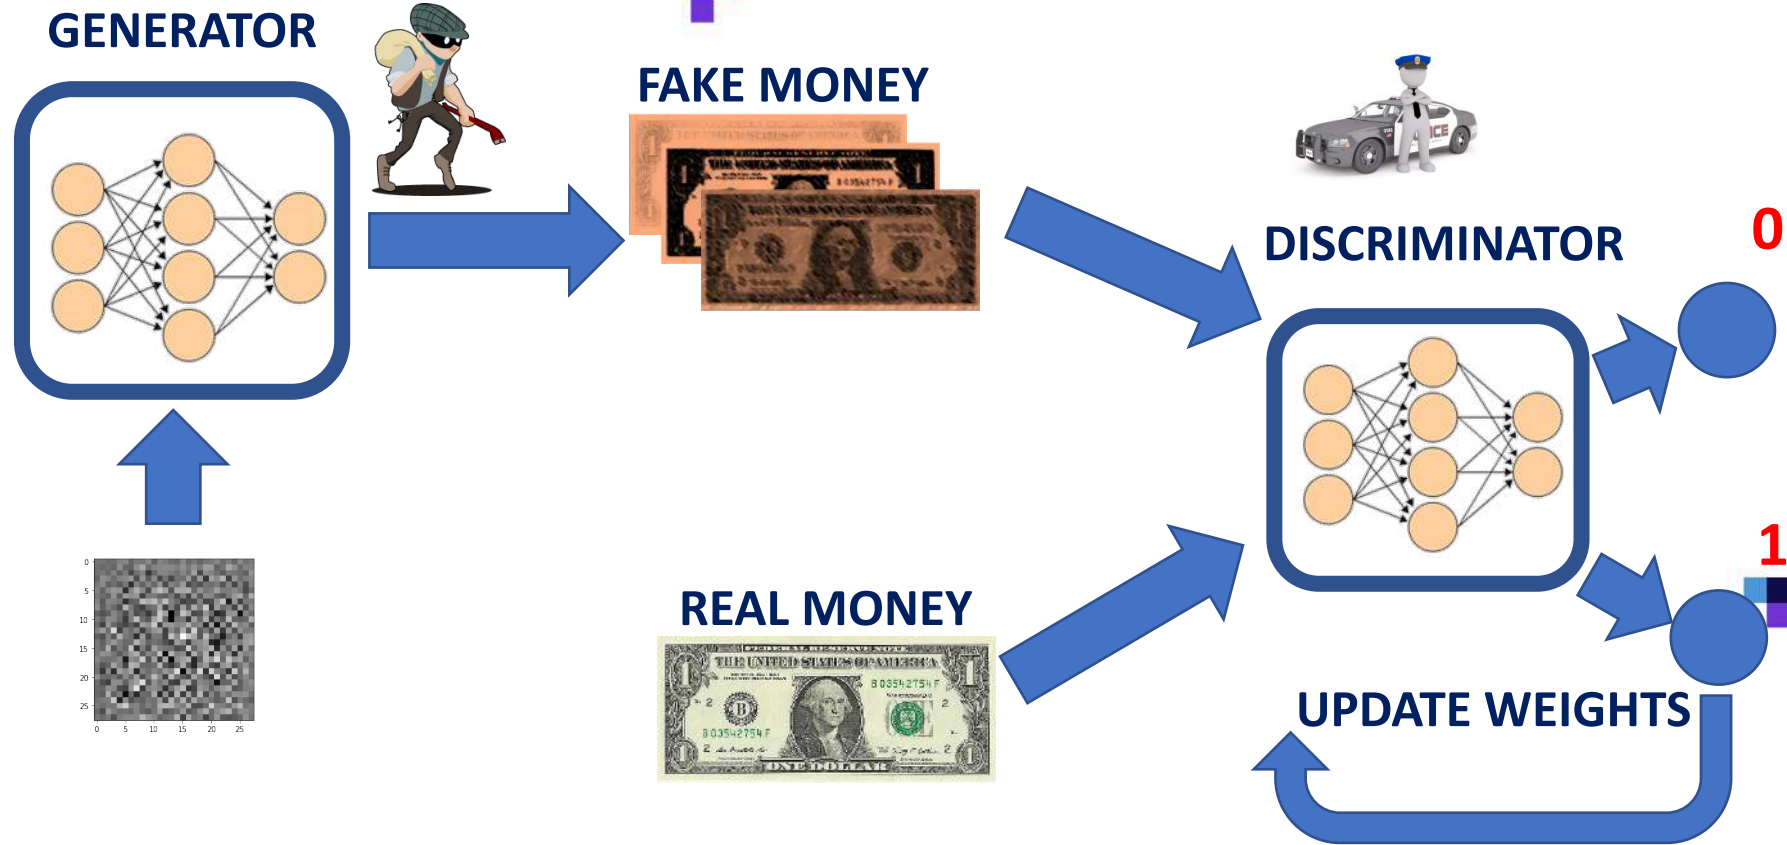


0

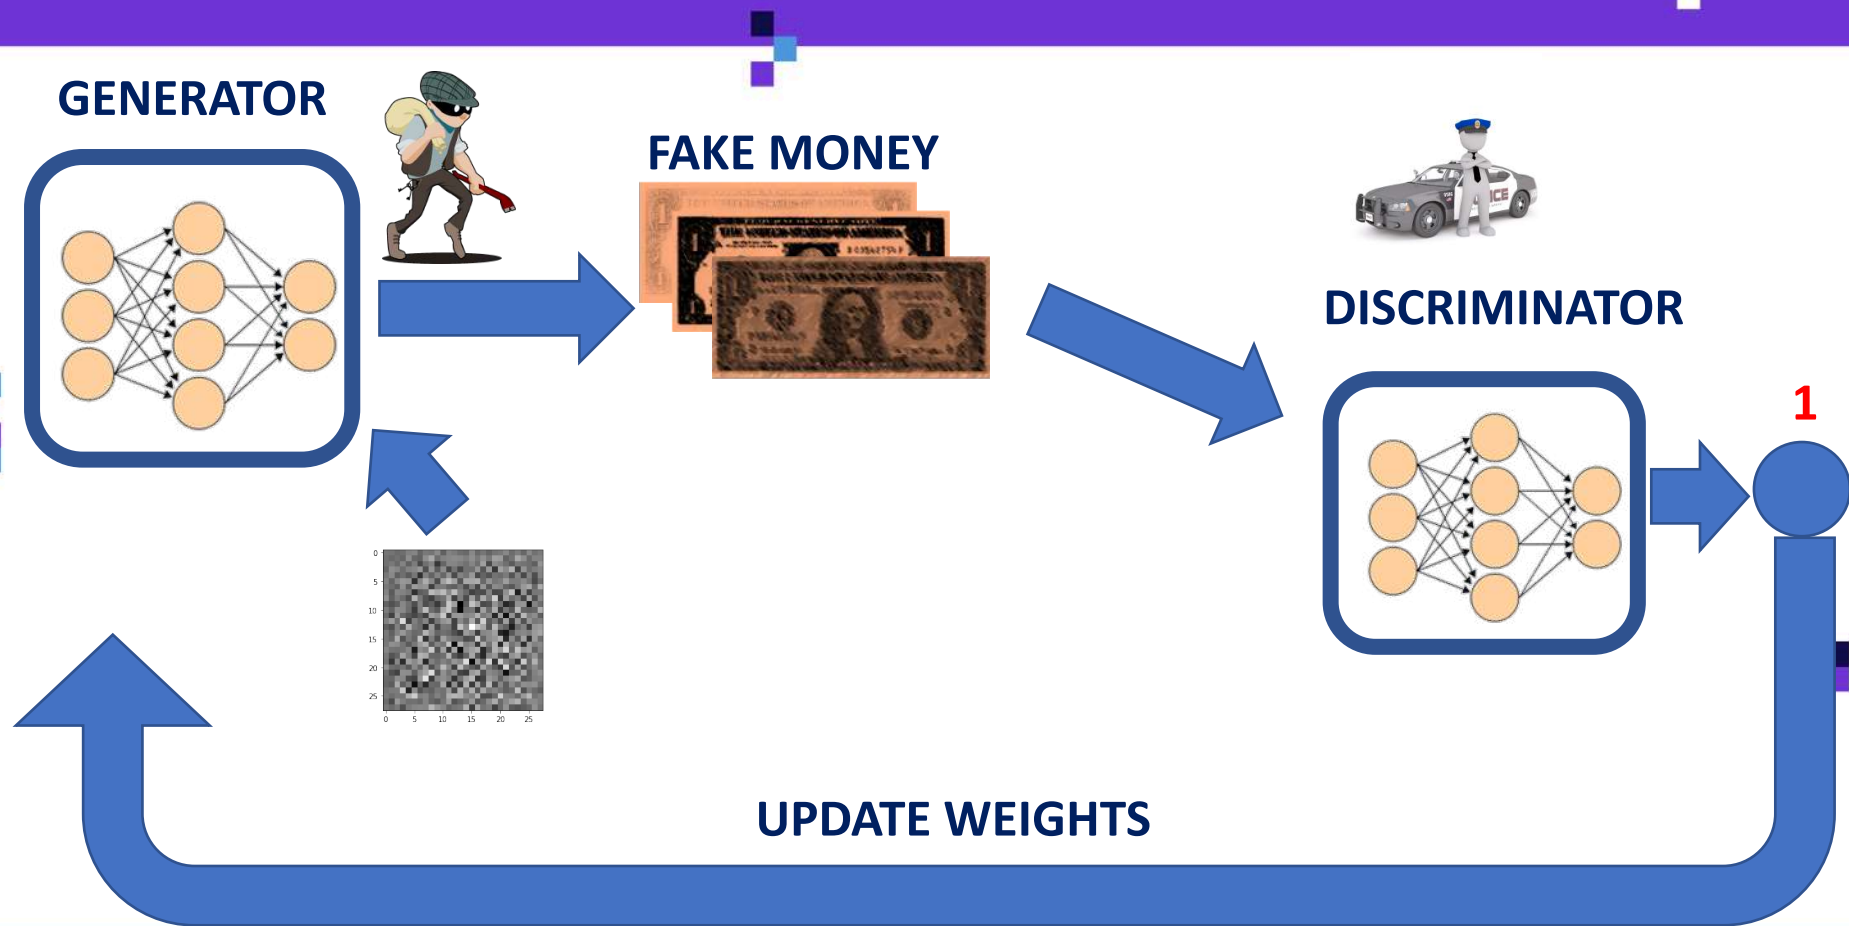
1



GANs – DISCRIMINATOR TRAINING




GANs – GENERATOR TRAINING



PLAN OF ATTACK – IMAGE SEGMENTATION

1. Intuition about segmentation
2. Implementation




Source of images: data-flair.training & deepsense.ai


INSTANCE VS. SEMANTIC SEGMENTATION




(a) Image classification



(b) Object localization



(c) Semantic segmentation



(d) Instance segmentation

Source: [Garcia et al](#)




Source: [NVIDIA Developer Blog](#)



Source: [Pexels](#)

MASK R-CNN


- It was built based on the previous object detection works: R-CNN (2013), Fast R-CNN (2015), and Faster R-CNN (2015), all by Girshick et al.



Source: <https://arxiv.org/abs/1703.06870>


R-CNN, FAST R-CNNs AND FASTER R-CNN

Original R-CNN Architecture




Source: [Girshick et al., 2013](#))

Fast R-CNN Architecture



Source: [Girshick et al., 2015](#))

Faster R-CNN Architecture



Source: [Girshick et al., 2015](#))